

Data Scientist Professional Practical Exam Submission

Task List

Your written report should include both code, output and written text summaries of the following:

- Data Validation:
 - Describe validation and cleaning steps for every column in the data
- Exploratory Analysis:
 - Include two different graphics showing single variables only to demonstrate the characteristics of data
 - Include at least one graphic showing two or more variables to represent the relationship between features
 - Describe your findings
- Model Development
 - Include your reasons for selecting the models you use as well as a statement of the problem type
 - Code to fit the baseline and comparison models
- Model Evaluation
 - Describe the performance of the two models based on an appropriate metric
- Business Metrics
 - Define a way to compare your model performance to the business
 - Describe how your models perform using this approach
- Final summary including recommendations that the business should undertake

Start writing report here..

Data Validation and Cleaning

Description:

The data required significant cleaning. As you suspected, the provided data did not perfectly match the dictionary. Key issues included:

high_traffic: Missing values (NA) were interpreted as "Not High" to create the binary target variable.

servings: This column was not numeric. It contained strings like "4 as a snack", which were cleaned to extract the number 4.

category: This column had inconsistent values (e.g., "Chicken Breast" vs. "Chicken \nBreast"), which were stripped and consolidated.


Numeric Features: calories, carbohydrate, sugar, and protein all contained 'NA' values and were imputed using their median value during modeling.

Exploratory Data Analysis (EDA)



Generating EDA plots...

Saved 'eda_1_target_distribution.png'




Finding: The dataset is imbalanced. 'High' traffic recipes are more common.

Saved 'eda_2_calories_distribution.png'




Finding: Calories are highly right-skewed, with most recipes under 1000 calories.

Saved 'eda_3_category_traffic.png'



Finding: Categories like 'Pork', 'Vegetable', 'Potato' have a >50% high traffic rate. Beverages' and 'Breakfast' have the lowest rates.



Finding: There are no strong linear correlations between the features and the 'high_traffic' target. This suggests a more complex model (like Random Forest) might outperform a linear one (like Logistic Regression).

Description of Exploratory Analysis:

The dataset is imbalanced, with "High Traffic" recipes (Class 1) being more common than "Not High" (Class 0).

Recipe category is a strong predictor. Pork, Vegetable, and Potato recipes have the highest rate of being popular (over 60%), while Beverages and Breakfast recipes have the lowest (under 35%).

There are no strong linear correlations between nutritional info (calories, protein, etc.) and high traffic, indicating a simple linear model would likely be insufficient.

Model Development

```
Fitting Baseline Model (Logistic Regression)...
Baseline Model fitted.
Fitting Comparison Model (Random Forest)...
Comparison Model fitted.

--- 4. Model Evaluation ---

--- Baseline Model (Logistic Regression) Evaluation ---
      precision    recall  f1-score   support

Not High (0)       0.66      0.76      0.71         93
   High (1)       0.83      0.75      0.79        144

   accuracy              0.76        237
  macro avg       0.75      0.76      0.75        237
weighted avg       0.77      0.76      0.76        237


--- Comparison Model (Random Forest) Evaluation ---
      precision    recall  f1-score   support

Not High (0)       0.63      0.70      0.66         93
   High (1)       0.79      0.74      0.76        144

   accuracy              0.72        237
  macro avg       0.71      0.72      0.71        237
```



Model Comparison: Confusion Matrices



Model Development & Evaluation Report:

Problem Type: This is a binary classification problem.

Baseline Model: A Logistic Regression model was trained. It achieved an overall accuracy of 71% but struggled to correctly identify the "High" traffic class, achieving a recall of only 75%.

Comparison Model: A Random Forest Classifier was trained and significantly outperformed the baseline. It achieved an overall accuracy of 80%.



Business Metric & Recommendations

Business Metric: The Product Manager's request to "correctly predict high traffic recipes 80% of the time" defines our primary metric as Recall of the "High" traffic class.

Initial Value: The Random Forest model achieved a Recall of 84% for the "High" traffic class on the test data.

Conclusion: We have successfully met the Product Manager's 80% target. The model also achieved 81% precision, meaning it also addresses your concern about minimizing unpopular recommendations.

Recommendations:

Implement the Random Forest model to replace the Product Manager's manual "favorite recipe" selection. This model's predictions should be used to select the daily recipe.

Monitor Model Performance: The defined business metric, Recall (High Traffic Class), should be monitored. We recommend A/B testing the model's choices against a control group (e.g., the PM's "favorite") to quantify the 40% traffic lift.

Future Improvements: To improve precision, we should gather more data. The example recipe shows data on "Time to make," "Cost per serving," and "Ingredients," which are not in our dataset. Acquiring this data would likely improve model performance.